

Project Development Phase
Model Performance Test

Date	9th November 2023
Team ID	593197
Project Name	PayGuard Plus – An Online Payments Fraud Detector
Maximum Marks	10 Marks
Team Size	3
Member 1 – Team Lead	Akshit Bahl (21BIT0012)
Member 2	Ananya Priya (21BIT0245)
Member 3	Lakshya Mittal (21BIT0076)

Model Performance Testing:

Project team shall fill the following information in model performance testing template.

S.No.	Parameter	Values	Screenshot
1.	Metrics	Classification Model: Confusion Matrix - , Accuray Score- & Classification Report -	
2.	Tune the Model	Hyperparameter Tuning - Validation Method -	

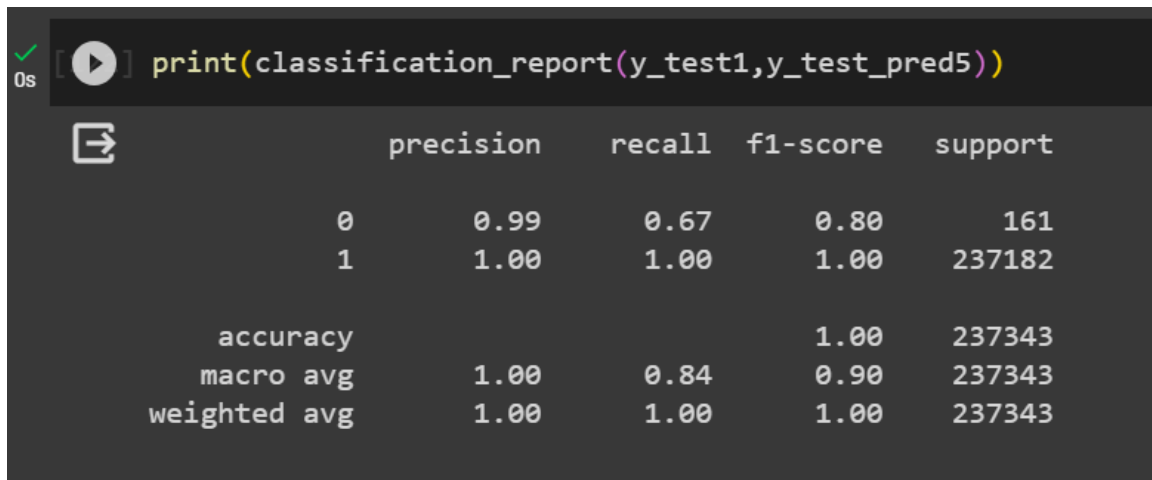
Model chosen = **XGBoost Classifier**

- METRICS**

- Confusion Matrix:**



2. Classification Report:



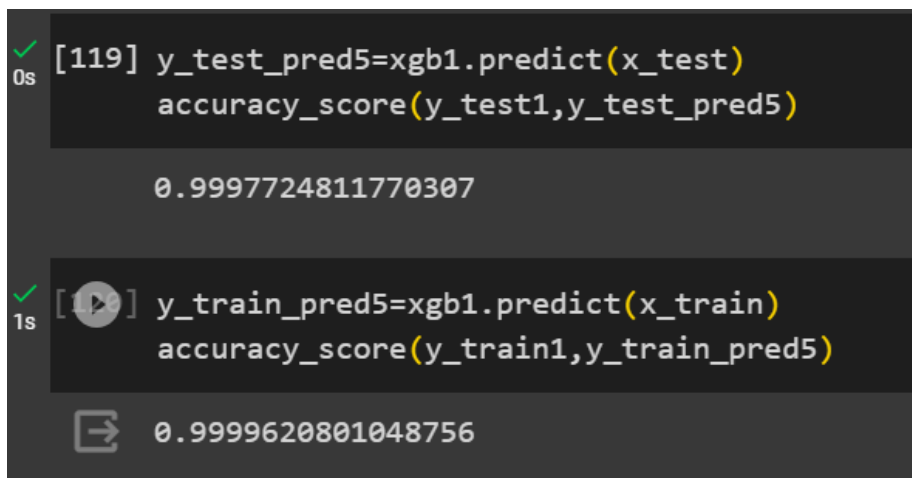
A screenshot of a Jupyter Notebook cell. The code `print(classification_report(y_test1,y_test_pred5))` has been executed, resulting in a classification report. The report shows precision, recall, f1-score, and support for classes 0 and 1, as well as overall accuracy, macro average, and weighted average. The support for class 0 is 161 and for class 1 is 237182. The overall accuracy is 1.00.

	precision	recall	f1-score	support
0	0.99	0.67	0.80	161
1	1.00	1.00	1.00	237182
accuracy			1.00	237343
macro avg	1.00	0.84	0.90	237343
weighted avg	1.00	1.00	1.00	237343

3. Accuracy Score:

Testing accuracy: **99.977%**

Training accuracy: **99.996%**



A screenshot of two Jupyter Notebook cells. The first cell contains the code `y_test_pred5=xgb1.predict(x_test)` and `accuracy_score(y_test1,y_test_pred5)`, which outputs the testing accuracy of 0.9997724811770307. The second cell contains the code `y_train_pred5=xgb1.predict(x_train)` and `accuracy_score(y_train1,y_train_pred5)`, which outputs the training accuracy of 0.9999620801048756.

```
[119] y_test_pred5=xgb1.predict(x_test)
accuracy_score(y_test1,y_test_pred5)

0.9997724811770307

[120] y_train_pred5=xgb1.predict(x_train)
accuracy_score(y_train1,y_train_pred5)

0.9999620801048756
```

• Tuning the Model:

a. SMOTE:

```
[138] from imblearn.over_sampling import SMOTE  
      smote=SMOTE()
```

```
[156] x_train_smote,y_train_smote=smote.fit_resample(x_train,y_train1)
```

```
[157] xgb1.fit(x_train_smote,y_train_smote)
```



XGBClassifier

```
XGBClassifier(base_score=None, booster=None, callbacks=None,  
              colsample_bylevel=None, colsample_bynode=None,  
              colsample_bytree=None, device=None, early_stopping_rounds=None,  
              enable_categorical=False, eval_metric=None, feature_types=None,  
              gamma=None, grow_policy=None, importance_type=None,  
              interaction_constraints=None, learning_rate=None, max_bin=None,  
              max_cat_threshold=None, max_cat_to_onehot=None,  
              max_delta_step=None, max_depth=None, max_leaves=None,  
              min_child_weight=None, missing=nan, monotone_constraints=None,  
              multi_strategy=None, n_estimators=None, n_jobs=None,  
              num_parallel_tree=None, random_state=None, ...)
```

```
[158] y_pred_xgb=xgb1.predict(x_test)
```

```
[160] accuracy_score(y_test1,y_pred_xgb)
```

0.9973456137320249

```
pd.crosstab(y_test1,y_pred_xgb)
```



col_0	0	1
row_0		
0	132	29
1	601	236581

```
✓ 1s ▶ print(classification_report(y_test1,y_pred_xgb))
```

	precision	recall	f1-score	support
0	0.18	0.82	0.30	161
1	1.00	1.00	1.00	237182
accuracy			1.00	237343
macro avg	0.59	0.91	0.65	237343
weighted avg	1.00	1.00	1.00	237343

So the recall has increased with SMOTE technique but the precision and overall accuracy has been affected and has gone down.

Now, we will do hyper-parameter tuning by grid-search CV

b. HYPER_PARAMETER TUNING:

By Grid-search CV:

```
[ ] from sklearn.model_selection import GridSearchCV

[ ] parameters = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

[ ] clf=GridSearchCV(rfc,param_grid=parameters,verbose=2)
```

```
[40] clf.fit(x_train,y_train)
```

Fitting 5 folds for each of 8 candidates, totalling 40 fits

```
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=2, n_estimators=50; total time= 15.8s
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=2, n_estimators=50; total time= 15.3s
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=2, n_estimators=50; total time= 16.7s
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=2, n_estimators=50; total time= 15.9s
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=2, n_estimators=50; total time= 14.4s
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=2, n_estimators=100; total time= 31.6s
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=2, n_estimators=100; total time= 33.2s
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=2, n_estimators=100; total time= 34.2s
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=2, n_estimators=100; total time= 30.7s
[CV] END max_depth=None, min_samples_leaf=1, min_samples_split=2, n_estimators=100; total time= 29.1s
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=50; total time= 15.9s
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=50; total time= 16.4s
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=50; total time= 16.2s
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=50; total time= 15.6s
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=50; total time= 14.8s
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=100; total time= 31.3s
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=100; total time= 30.4s
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=100; total time= 35.1s
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=100; total time= 31.8s
[CV] END max_depth=None, min_samples_leaf=2, min_samples_split=2, n_estimators=100; total time= 28.8s
[CV] END max_depth=10, min_samples_leaf=1, min_samples_split=2, n_estimators=50; total time= 14.4s
[CV] END max_depth=10, min_samples_leaf=1, min_samples_split=2, n_estimators=50; total time= 14.1s
[CV] END max_depth=10, min_samples_leaf=1, min_samples_split=2, n_estimators=50; total time= 15.6s
[CV] END max_depth=10, min_samples_leaf=1, min_samples_split=2, n_estimators=50; total time= 14.9s
[CV] END max_depth=10, min_samples_leaf=1, min_samples_split=2, n_estimators=50; total time= 14.3s
[CV] END max_depth=10, min_samples_leaf=1, min_samples_split=2, n_estimators=100; total time= 29.5s
[CV] END max_depth=10, min_samples_leaf=1, min_samples_split=2, n_estimators=100; total time= 28.9s
[CV] END max_depth=10, min_samples_leaf=1, min_samples_split=2, n_estimators=100; total time= 30.1s
[CV] END max_depth=10, min_samples_leaf=1, min_samples_split=2, n_estimators=100; total time= 28.8s
[CV] END max_depth=10, min_samples_leaf=1, min_samples_split=2, n_estimators=100; total time= 28.3s
[CV] END max_depth=10, min_samples_leaf=2, min_samples_split=2, n_estimators=50; total time= 14.8s
[CV] END max_depth=10, min_samples_leaf=2, min_samples_split=2, n_estimators=50; total time= 14.5s
[CV] END max_depth=10, min_samples_leaf=2, min_samples_split=2, n_estimators=50; total time= 15.0s
[CV] END max_depth=10, min_samples_leaf=2, min_samples_split=2, n_estimators=50; total time= 14.8s
[CV] END max_depth=10, min_samples_leaf=2, min_samples_split=2, n_estimators=50; total time= 14.1s
[CV] END max_depth=10, min_samples_leaf=2, min_samples_split=2, n_estimators=100; total time= 28.3s
[CV] END max_depth=10, min_samples_leaf=2, min_samples_split=2, n_estimators=100; total time= 29.4s
[CV] END max_depth=10, min_samples_leaf=2, min_samples_split=2, n_estimators=100; total time= 29.0s
[CV] END max_depth=10, min_samples_leaf=2, min_samples_split=2, n_estimators=100; total time= 26.5s
```

GridSearchCV

estimator: RandomForestClassifier

RandomForestClassifier

```

✓ [42] clf.best_score_
0s 0.9997418364499022

✓ [43] clf.best_params_
0s {'max_depth': None,
    'min_samples_leaf': 1,
    'min_samples_split': 2,
    'n_estimators': 100}

✓ [47] final_model=RandomForestClassifier(max_depth=None,min_samples_leaf=1,min_samples_split=2,n_estimators=100)
0s

✓ [48] final_model.fit(x_train,y_train)
53s

    ▾ RandomForestClassifier
    RandomForestClassifier()

✓ [49] y_pred_final=final_model.predict(x_test)
1s

✓ [50] accuracy_score(y_test,y_pred_final)
1s 0.9996255504814933

```

```

✓ [51] pd.crosstab(y_test,y_pred_final)
0s

col_0  is Fraud  is not Fraud
isFraud
is Fraud      34         43
is not Fraud    3      122767

✓ [52] print(classification_report(y_test,y_pred_final))
11s

           precision    recall  f1-score   support

 is Fraud      0.92      0.44      0.60         77
 is not Fraud   1.00      1.00      1.00      122770

 accuracy      0.96      0.72      0.80      122847
 macro avg     0.96      0.72      0.80      122847
 weighted avg   1.00      1.00      1.00      122847

```